

## Analyzing 9-1-1 Call Data using Bayesian Regression

This is an analysis of Poisson and negative binomial regression models using both Bayesian and frequentist methods. The data is aggregated data from a Whitcom 9-1-1 public records request as well as additional information regarding home football games for Washington State University. Of the five columns in the data set (Count, Date, Hour, Month, Game.Day), all but Date are of interest for this analysis. Using Count as the response, we will fit four models using a combination of Hour, Month, and Game.Day. All in all, there will be sixteen models evaluated.

```
library(rjags)

## Loading required package: coda


## Linked to JAGS 4.3.0

## Loaded modules: basemod,bugs

library(ggplot2)
library(ggpubr)
library(MASS)

# Read in the data
data = read.csv('Data/Project_data.csv', header=TRUE)
data$Hour = as.factor(data$Hour)
data$Month = as.factor(data$Month)
data$Game.Day = as.factor(data$Game.Day)

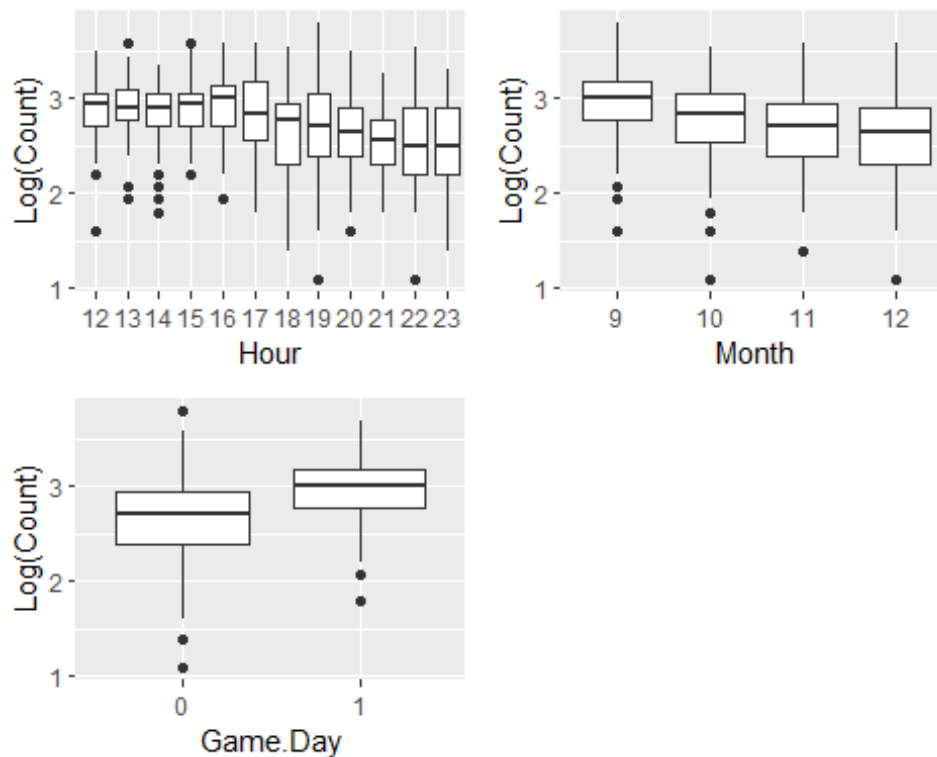
# Histogram of Count
ggplot(data, aes(x=Count))+geom_histogram(bins=44)+
  ylab('Frequency of Count')+ggtitle('Histogram of Count')
```



```
# Plot Log(Count) against each predictor
ggarrange(
  ggplot(data, aes(x=Hour, y=log(Count)))+geom_boxplot()+
    xlab('Hour')+ylab('Log(Count)'),

  ggplot(data, aes(x=Month, y=log(Count)))+geom_boxplot()+
    xlab('Month')+ylab('Log(Count)'),

  ggplot(data, aes(x=Game.Day, y=log(Count)))+geom_boxplot()+
    xlab('Game.Day')+ylab('Log(Count)')
```



It appears there may exist predictive relationships between the log of Count and each predictor.

```
# Create the response matrix
```

```
Y = data$Count
```

```
# Create the design matrix
```

```
X = model.matrix(Count~Hour+Month+Game.Day, data=data)
```

```
head(X)
```

```
##      (Intercept) Hour13 Hour14 Hour15 Hour16 Hour17 Hour18 Hour19 Hour20
Hour21
## 1             1      0      0      0      0      0      0      0      0
0
## 2             1      1      0      0      0      0      0      0      0
0
## 3             1      0      1      0      0      0      0      0      0
0
## 4             1      0      0      1      0      0      0      0      0
0
## 5             1      0      0      0      1      0      0      0      0
0
## 6             1      0      0      0      0      1      0      0      0
0
##      Hour22 Hour23 Month10 Month11 Month12 Game.Day1
## 1         0      0      0      0      0         1
## 2         0      0      0      0      0         1
```

```
## 3      0      0      0      0      0      1
## 4      0      0      0      0      0      1
## 5      0      0      0      0      0      1
## 6      0      0      0      0      0      1
```

```
set.seed(7)
```

```
# 80/20 training/testing split
train_rows = sample(nrow(data), 0.8*nrow(data))
```

```
# Matrix form
train_X = X[train_rows,]
test_X = X[-train_rows,]
```

```
train_Y = Y[train_rows]
test_Y = Y[-train_rows]
```

```
# Dataframe form
train_data = data[train_rows,]
test_data = data[-train_rows,]
```

```
# Poisson models and priors
p_jags = '
model{
  for(i in 1:n){

    # LIKELIHOODS

    # Model 1 (Count~Hour)
    Y1[i] ~ dpois(lambda1[i])
    log(lambda1[i]) <- inprod(X1[i,], beta1[])

    # Model 2 (Count~Hour+Month)
    Y2[i] ~ dpois(lambda2[i])
    log(lambda2[i]) <- inprod(X2[i,], beta2[])

    # Model 3 (Count~Hour+Game.Day)
    Y3[i] ~ dpois(lambda3[i])
    log(lambda3[i]) <- inprod(X3[i,], beta3[])

    # Model 4 (Count~Hour+Month+Game.Day)
    Y4[i] ~ dpois(lambda4[i])
    log(lambda4[i]) <- inprod(X4[i,], beta4[])
  }
}
```

```
# PRIORS
```

```

# Priors for Model 1
for(p1 in 1:12){
  beta1[p1] ~ dnorm(0, 0.001)
}

# Priors for Model 2
for(p2 in 1:15){
  beta2[p2] ~ dnorm(0, 0.001)
}

# Priors for Model 3
for(p3 in 1:13){
  beta3[p3] ~ dnorm(0, 0.001)
}

# Priors for Model 4
for(p4 in 1:16){
  beta4[p4] ~ dnorm(0, 0.001)
}

# PREDICTIONS USING THE TEST SET

for(i in 1:m){

  # Model 1 Predictions
  log(lambda1_star[i]) <- inprod(X1_test[i,], beta1[])
  pred1[i] ~ dpois(lambda1_star[i])

  # Model 2 Predictions
  log(lambda2_star[i]) <- inprod(X2_test[i,], beta2[])
  pred2[i] ~ dpois(lambda2_star[i])

  # Model 3 Predictions
  log(lambda3_star[i]) <- inprod(X3_test[i,], beta3[])
  pred3[i] ~ dpois(lambda3_star[i])

  # Model 4 Predictions
  log(lambda4_star[i]) <- inprod(X4_test[i,], beta4[])
  pred4[i] ~ dpois(lambda4_star[i])
}
}

# Data
p_data = list(n=nrow(train_X),
              m=nrow(test_X),

```

```

X1=train_X[,1:12],
X2=train_X[,1:15],
X3=train_X[,c(1:12, 16)],
X4=train_X,

X1_test=test_X[,1:12],
X2_test=test_X[,1:15],
X3_test=test_X[,c(1:12, 16)],
X4_test=test_X,

Y1=train_Y,
Y2=train_Y,
Y3=train_Y,
Y4=train_Y)

# Initialize models
p_models = jags.model(file=textConnection(p_jags), data=p_data,
                      inits=list(.RNG.name = 'base:Wichmann-Hill',
                                .RNG.seed = 7))

## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 2032
##   Unobserved stochastic nodes: 568
##   Total graph size: 41104
##
## Initializing model

# Burn in
update(p_models, n.iter=1000)

# Number of iterations
iter = 1e4

# Sample
p_outputs = coda.samples(p_models,
                         variable.names = c('beta1', 'beta2',
                                             'beta3', 'beta4',
                                             'lambda1', 'lambda2',
                                             'lambda3', 'lambda4',
                                             'pred1', 'pred2',
                                             'pred3', 'pred4'),
                         n.iter = iter)

# Create 4 frequentist models
p1 = glm(Count~Hour, data=train_data, family='poisson')
p2 = glm(Count~Hour+Month, data=train_data, family='poisson')
p3 = glm(Count~Hour+Game.Day, data=train_data, family='poisson')

```

```

p4 = glm(Count~Hour+Month+Game.Day, data=train_data, family='poisson')

# Extract the estimate of each Bayesian node
p_est = summary(p_outputs)$statistics[,1]

# Extract the 95% equi-tailed credible set of each node
p_quants = summary(p_outputs)$quantiles[,c(1,5)]

# Create a table of Bayesian coefficients with their credible set
p_beta1 = cbind(Estimate=p_est[1:12], p_quants[1:12,])

p_beta2 = cbind(Estimate=p_est[13:27], p_quants[13:27,])

p_beta3 = cbind(Estimate=p_est[28:40], p_quants[28:40,])

p_beta4 = cbind(p_est[41:56], p_quants[41:56,])

# Rename row names for better interpretation
rownames(p_beta1) = c('(Intercept)', 'Hour13', 'Hour14', 'Hour15', 'Hour16',
                      'Hour17', 'Hour18', 'Hour19', 'Hour20', 'Hour21',
                      'Hour22', 'Hour23')

rownames(p_beta2) = c('(Intercept)', 'Hour13', 'Hour14', 'Hour15', 'Hour16',
                      'Hour17', 'Hour18', 'Hour19', 'Hour20', 'Hour21',
                      'Hour22', 'Hour23', 'Month10', 'Month11', 'Month12')

rownames(p_beta3) = c('(Intercept)', 'Hour13', 'Hour14', 'Hour15', 'Hour16',
                      'Hour17', 'Hour18', 'Hour19', 'Hour20', 'Hour21',
                      'Hour22', 'Hour23', 'Game.Day')

rownames(p_beta4) = c('(Intercept)', 'Hour13', 'Hour14', 'Hour15', 'Hour16',
                      'Hour17', 'Hour18', 'Hour19', 'Hour20', 'Hour21',
                      'Hour22', 'Hour23', 'Month10', 'Month11', 'Month12',
                      'Game.Day')

# Print Bayesian coefficients with their credible sets
p_beta1

```

	Estimate	2.5%	97.5%
(Intercept)	2.942620851	2.87257437	3.011388907
Hour13	-0.004945450	-0.09936393	0.089815672
Hour14	-0.089667594	-0.19371536	0.016635046
Hour15	0.007390016	-0.09026404	0.106847978
Hour16	0.024506690	-0.07092365	0.118258916
Hour17	-0.034642500	-0.13233598	0.062635910
Hour18	-0.204641979	-0.30987376	-0.103441382
Hour19	-0.113044809	-0.21903480	-0.009888539

## Hour20	-0.309521262	-0.41665880	-0.201518591
## Hour21	-0.309227028	-0.41514291	-0.202240589
## Hour22	-0.292859235	-0.39787854	-0.188053597
## Hour23	-0.287710008	-0.39335266	-0.181559895

#### p\_beta2

##	Estimate	2.5%	97.5%
## (Intercept)	3.101231362	3.02509989	3.17309777
## Hour13	-0.008741151	-0.10313219	0.08612208
## Hour14	-0.081526226	-0.18100219	0.01945024
## Hour15	0.005592348	-0.08887491	0.09990170
## Hour16	0.018073521	-0.07385414	0.11032883
## Hour17	-0.035113494	-0.12860395	0.06213752
## Hour18	-0.215515432	-0.31607777	-0.11613118
## Hour19	-0.132089785	-0.23504108	-0.02676695
## Hour20	-0.292864970	-0.39772239	-0.18597796
## Hour21	-0.308105770	-0.40877779	-0.20742150
## Hour22	-0.296172721	-0.40058463	-0.19290069
## Hour23	-0.298484525	-0.40172863	-0.19371928
## Month10	-0.154200952	-0.21246416	-0.09624511
## Month11	-0.226754637	-0.28679946	-0.16771271
## Month12	-0.262667827	-0.31927433	-0.20613843

#### p\_beta3

##	Estimate	2.5%	97.5%
## (Intercept)	2.858590629	2.78648098	2.92726859
## Hour13	-0.017652015	-0.11184657	0.07809581
## Hour14	-0.086580171	-0.18965667	0.01904538
## Hour15	-0.002919066	-0.09931943	0.09697755
## Hour16	0.023880498	-0.06646064	0.11898009
## Hour17	-0.041752821	-0.13994982	0.05601537
## Hour18	-0.217268864	-0.31872138	-0.11327866
## Hour19	-0.136501112	-0.23910823	-0.03116623
## Hour20	-0.295711422	-0.40386500	-0.18755878
## Hour21	-0.317991644	-0.41968701	-0.21243387
## Hour22	-0.298609156	-0.40347264	-0.19254296
## Hour23	-0.300589512	-0.40546489	-0.19560144
## Game.Day	0.272029585	0.22749658	0.31548137

#### p\_beta4

##	Estimate	2.5%	97.5%
## (Intercept)	2.957794171	2.87388225	3.03579226
## Hour13	-0.014139694	-0.10448787	0.08160450
## Hour14	-0.078991353	-0.18388751	0.02625790
## Hour15	-0.003598471	-0.09814970	0.09305833
## Hour16	0.018594269	-0.07319965	0.11615175
## Hour17	-0.040167645	-0.13710844	0.06042960
## Hour18	-0.222317174	-0.32201503	-0.11918210



```
## Hour19      -0.145483010 -0.24914736 -0.04456522
## Hour20      -0.284629760 -0.39228944 -0.17557722
## Hour21      -0.314237260 -0.41671255 -0.20947911
## Hour22      -0.296995292 -0.39973330 -0.19112992
## Hour23      -0.297289228 -0.40295097 -0.19378172
## Month10     -0.074213358 -0.13455659 -0.01408321
## Month11     -0.184763643 -0.24516446 -0.12675892
## Month12     -0.116067482 -0.18157043 -0.05158707
## Game.Day    0.242713607  0.19243433  0.29374964
```

```
# Print GLM coefficients with 95% confidence intervals
cbind(Estimate=coef(p1), confint(p1))
```

```
## Waiting for profiling to be done...
```

```
##           Estimate      2.5 %      97.5 %
## (Intercept)  2.944438979  2.87665118  3.01072661
## Hour13      -0.006741599 -0.10066686  0.08723778
## Hour14      -0.090768588 -0.19339493  0.01135996
## Hour15       0.006397974 -0.09060525  0.10326533
## Hour16       0.023743343 -0.06996906  0.11752279
## Hour17      -0.036141005 -0.13301160  0.06060202
## Hour18      -0.205852054 -0.30742715 -0.10470564
## Hour19      -0.114498960 -0.21870083 -0.01089971
## Hour20      -0.310891305 -0.41890089 -0.20374386
## Hour21      -0.310154928 -0.41350256 -0.20735290
## Hour22      -0.293820827 -0.39807423 -0.19017351
## Hour23      -0.288906812 -0.39300607 -0.18540330
```

```
cbind(Estimate=coef(p2), confint(p2))
```

```
## Waiting for profiling to be done...
```

```
##           Estimate      2.5 %      97.5 %
## (Intercept)  3.105192007  3.02975267  3.17929863
## Hour13      -0.011975309 -0.10593695  0.08204045
## Hour14      -0.083297357 -0.18593849  0.01884613
## Hour15       0.002609741 -0.09442406  0.09950773
## Hour16       0.015514684 -0.07823746  0.10933385
## Hour17      -0.037500893 -0.13437725  0.05924791
## Hour18      -0.217514065 -0.31913038 -0.11632607
## Hour19      -0.135002700 -0.23931993 -0.03128670
## Hour20      -0.294227755 -0.40234489 -0.18697096
## Hour21      -0.310891391 -0.41424679 -0.20808152
## Hour22      -0.298349915 -0.40261457 -0.19469120
## Hour23      -0.301237216 -0.40545151 -0.19761716
## Month10     -0.154842314 -0.21387030 -0.09596761
## Month11     -0.227366427 -0.28668220 -0.16820917
## Month12     -0.263446248 -0.32165845 -0.20534945
```

```
cbind(Estimate=coef(p3), confint(p3))
```

```
## Waiting for profiling to be done...

##           Estimate      2.5 %      97.5 %
## (Intercept)  2.857902288  2.78839812  2.92596220
## Hour13      -0.015387387 -0.10932408  0.07860342
## Hour14      -0.085411560 -0.18804196  0.01672109
## Hour15      -0.001657116 -0.09866995  0.09521987
## Hour16       0.025547101 -0.06816580  0.11932706
## Hour17      -0.039989316 -0.13686211  0.05675592
## Hour18      -0.216326106 -0.31791665 -0.11516411
## Hour19      -0.135011655 -0.23927069 -0.03135460
## Hour20      -0.294072942 -0.40212095 -0.18688643
## Hour21      -0.316511460 -0.41986471 -0.21370376
## Hour22      -0.297669139 -0.40192458 -0.19401975
## Hour23      -0.299380864 -0.40349519 -0.19586212
## Game.Day1    0.272233268  0.22789301  0.31639676
```

```
cbind(Estimate=coef(p4), confint(p4))
```

```
## Waiting for profiling to be done...

##           Estimate      2.5 %      97.5 %
## (Intercept)  2.960284045  2.87832701  3.04101719
## Hour13      -0.014149184 -0.10811264  0.07986838
## Hour14      -0.080471848 -0.18311646  0.02167515
## Hour15      -0.004086963 -0.10112715  0.09281747
## Hour16       0.017358979 -0.07639458  0.11117956
## Hour17      -0.040464876 -0.13734320  0.05628590
## Hour18      -0.222972212 -0.32459248 -0.12178022
## Hour19      -0.146068739 -0.25041352 -0.04232487
## Hour20      -0.285210604 -0.39335712 -0.17792384
## Hour21      -0.315453179 -0.41881236 -0.21263948
## Hour22      -0.297678522 -0.40194387 -0.19401910
## Hour23      -0.298586476 -0.40280198 -0.19496521
## Month10     -0.074867288 -0.13629049 -0.01358229
## Month11     -0.185543823 -0.24549656 -0.12574728
## Month12     -0.117472250 -0.18389697 -0.05099900
## Game.Day1    0.242193943  0.19125180  0.29311039
```

Poisson deviance residuals  $d_i = \text{sign}(Y_i - \lambda_i) \sqrt{2[Y_i \log(Y_i/\lambda_i) - (Y_i - \lambda_i)]}$

```
# Extract fitted values for Bayesian models
```

```
p_fv1 = p_est[57:564]
```

```
p_fv2 = p_est[565:1072]
```

```
p_fv3 = p_est[1073:1580]
```

```
p_fv4 = p_est[1581:2088]
```

```
# Calculate deviance residuals for Bayesian models
```

```
pois_dev_res = function(fv){
```

```
  dr = sign(train_Y-fv)*sqrt(2*(train_Y*log(train_Y/fv)-(train_Y-fv)))
```

```
  return(dr)
```

```

}

p_dr = data.frame(
  p_dr1 = pois_dev_res(p_fv1),
  p_dr2 = pois_dev_res(p_fv2),
  p_dr3 = pois_dev_res(p_fv3),
  p_dr4 = pois_dev_res(p_fv4),
  row.names = c(1:508))

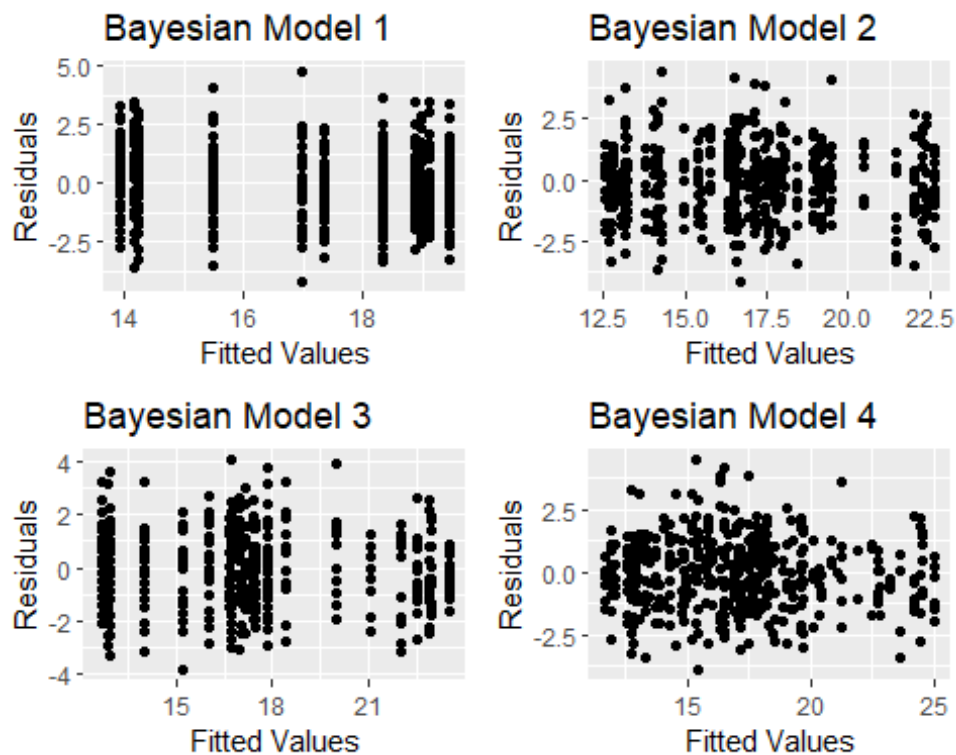
# Residual analysis
# Residuals vs fitted values
ggarrange(
  ggplot(p_dr, aes(x=p_fv1, y=p_dr1))+geom_point()+
    xlab('Fitted Values')+ylab('Residuals')+
    ggtitle('Bayesian Model 1'),

  ggplot(p_dr, aes(x=p_fv2, y=p_dr2))+geom_point()+
    xlab('Fitted Values')+ylab('Residuals')+
    ggtitle('Bayesian Model 2'),

  ggplot(p_dr, aes(x=p_fv3, y=p_dr3))+geom_point()+
    xlab('Fitted Values')+ylab('Residuals')+
    ggtitle('Bayesian Model 3'),

  ggplot(p_dr, aes(x=p_fv4, y=p_dr4))+geom_point()+
    xlab('Fitted Values')+ylab('Residuals')+
    ggtitle('Bayesian Model 4'))

```

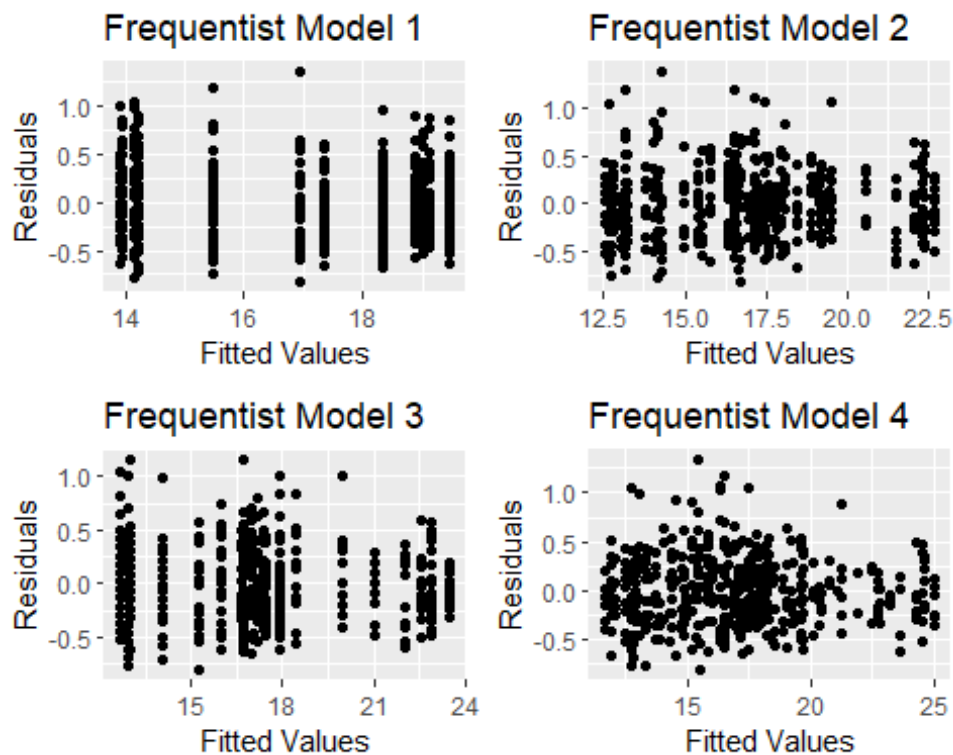


```
ggarrange(
  ggplot(p1, aes(p1$fitted.values, p1$residuals))+
    geom_point()+xlab('Fitted Values')+
    ylab('Residuals')+ggtitle('Frequentist Model 1'),

  ggplot(p2, aes(p2$fitted.values, p2$residuals))+
    geom_point()+xlab('Fitted Values')+
    ylab('Residuals')+ggtitle('Frequentist Model 2'),

  ggplot(p3, aes(p3$fitted.values, p3$residuals))+
    geom_point()+xlab('Fitted Values')+
    ylab('Residuals')+ggtitle('Frequentist Model 3'),

  ggplot(p4, aes(p4$fitted.values, p4$residuals))+
    geom_point()+xlab('Fitted Values')+
    ylab('Residuals')+ggtitle('Frequentist Model 4'))
```



*# Histogram of the Bayesian residuals*

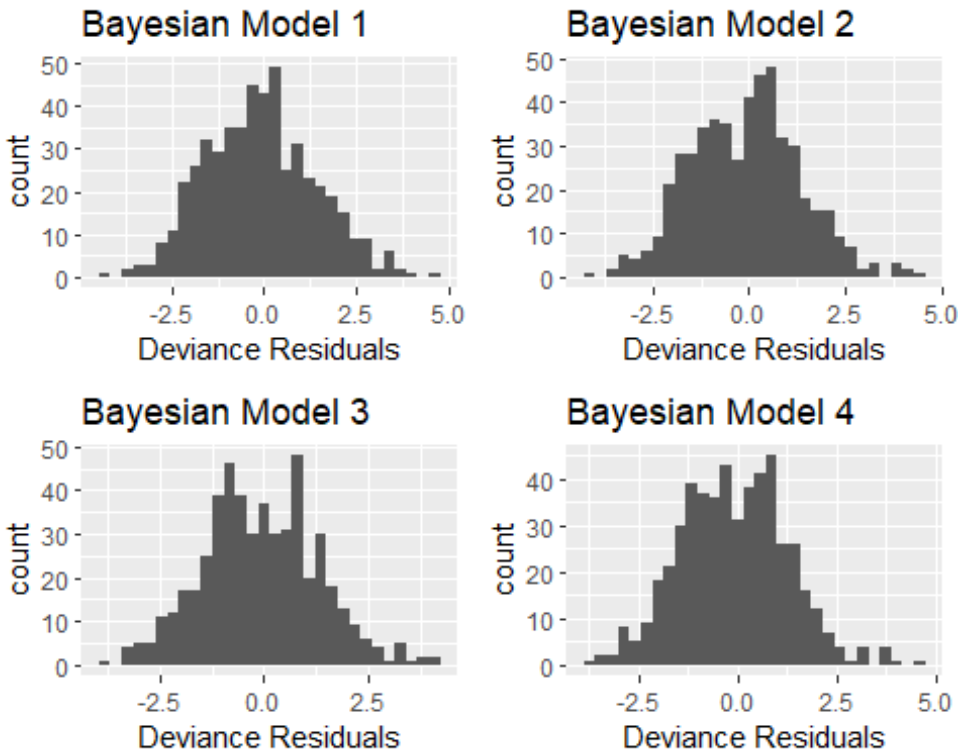
```
ggarrange(
  ggplot(p_dr, aes(x=p_dr1))+geom_histogram()+
    xlab('Deviance Residuals')+ggtitle('Bayesian Model 1'),

  ggplot(p_dr, aes(x=p_dr2))+geom_histogram()+
    xlab('Deviance Residuals')+ggtitle('Bayesian Model 2'),

  ggplot(p_dr, aes(x=p_dr3))+geom_histogram()+
    xlab('Deviance Residuals')+ggtitle('Bayesian Model 3'),

  ggplot(p_dr, aes(x=p_dr4))+geom_histogram()+
    xlab('Deviance Residuals')+ggtitle('Bayesian Model 4'))
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



*# Histogram of frequentist residuals*

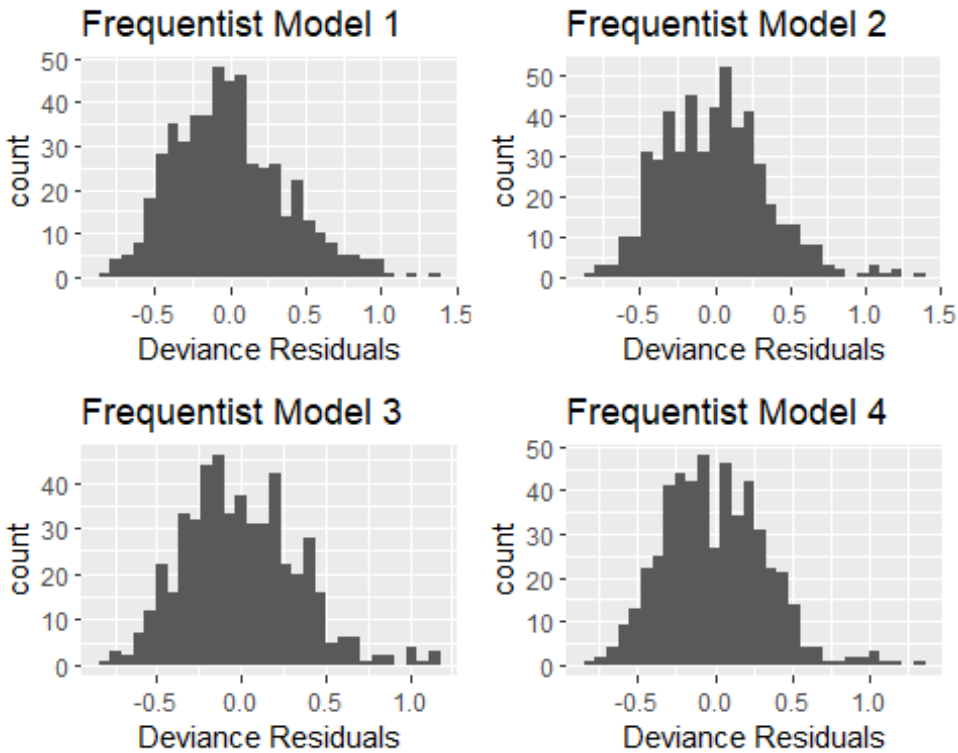
```
ggarrange(
  ggplot(p1, aes(x=p1$residuals))+geom_histogram()+
    xlab('Deviance Residuals')+ggtitle('Frequentist Model 1'),

  ggplot(p2, aes(x=p2$residuals))+geom_histogram()+
    xlab('Deviance Residuals')+ggtitle('Frequentist Model 2'),

  ggplot(p3, aes(x=p3$residuals))+geom_histogram()+
    xlab('Deviance Residuals')+ggtitle('Frequentist Model 3'),

  ggplot(p4, aes(x=p4$residuals))+geom_histogram()+
    xlab('Deviance Residuals')+ggtitle('Frequentist Model 4'))
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



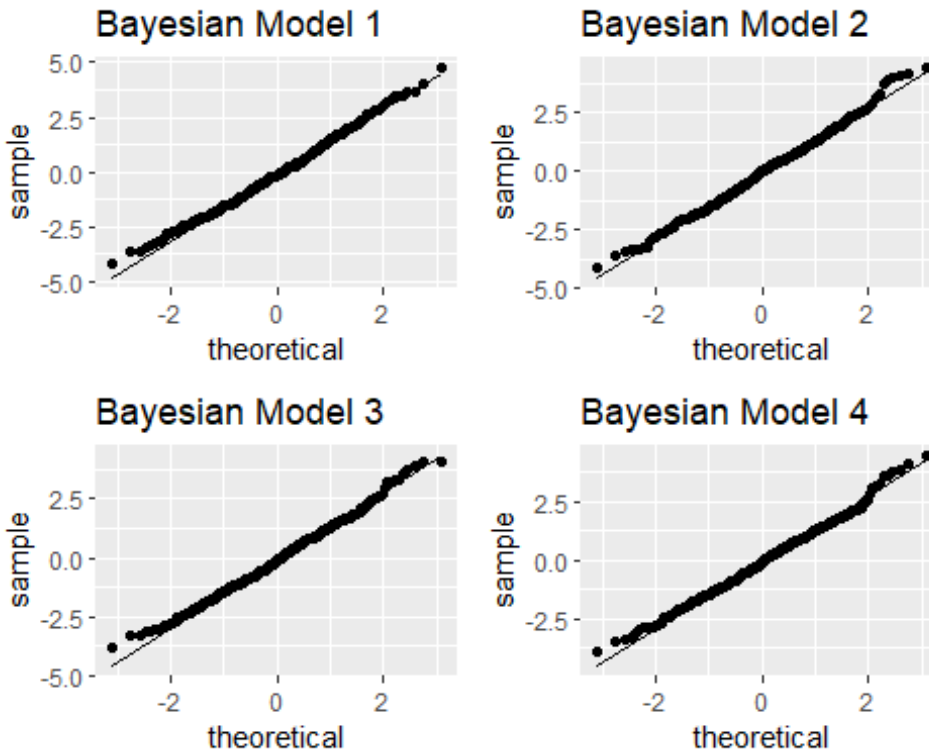
*# QQ plot of the Bayesian residuals*

```
ggarrange(ggplot(p_dr, aes(sample=p_dr1))+stat_qq()+stat_qq_line()+
  ggtitle('Bayesian Model 1'),

  ggplot(p_dr, aes(sample=p_dr2))+stat_qq()+stat_qq_line()+
  ggtitle('Bayesian Model 2'),

  ggplot(p_dr, aes(sample=p_dr3))+stat_qq()+stat_qq_line()+
  ggtitle('Bayesian Model 3'),

  ggplot(p_dr, aes(sample=p_dr4))+stat_qq()+stat_qq_line()+
  ggtitle('Bayesian Model 4'))
```




```
# Histogram of the frequentist residuals
ggarrange(ggplot(p1, aes(sample=p1$residuals))+stat_qq()+stat_qq_line()+
  ggtitle('Frequentist Model 1'),

  ggplot(p2, aes(sample=p2$residuals))+stat_qq()+stat_qq_line()+
  ggtitle('Frequentist Model 2'),

  ggplot(p3, aes(sample=p3$residuals))+stat_qq()+stat_qq_line()+
  ggtitle('Frequentist Model 3'),

  ggplot(p4, aes(sample=p4$residuals))+stat_qq()+stat_qq_line()+
  ggtitle('Frequentist Model 4'))
```





```
# Calculate the deviance of each Bayesian model
p_d1 = sum(p_dr$p_dr1^2)
p_d2 = sum(p_dr$p_dr2^2)
p_d3 = sum(p_dr$p_dr3^2)
p_d4 = sum(p_dr$p_dr4^2)

# Summary table (deviance, p-value, dispersion)
p_gof = data.frame(
  Deviance_B = c(p_d1, p_d2, p_d3, p_d4),

  GOF_B = c(1-pchisq(p_d1, summary(p1)$df.residual),
            1-pchisq(p_d2, summary(p2)$df.residual),
            1-pchisq(p_d3, summary(p3)$df.residual),
            1-pchisq(p_d4, summary(p4)$df.residual)),

  Dispersion_B = c(p_d1/summary(p1)$df.residual,
                   p_d2/summary(p1)$df.residual,
                   p_d3/summary(p1)$df.residual,
                   p_d4/summary(p1)$df.residual),

  Deviance_F = c(deviance(p1),
                 deviance(p2),
                 deviance(p3),
                 deviance(p4)),

  GOF_F = c(1-pchisq(deviance(p1), summary(p1)$df.residual),
```

```

1-pchisq(deviance(p2), summary(p2)$df.residual),
1-pchisq(deviance(p3), summary(p3)$df.residual),
1-pchisq(deviance(p4), summary(p4)$df.residual)),

Dispersion_B = c(deviance(p1)/summary(p1)$df.residual,
                  deviance(p2)/summary(p2)$df.residual,
                  deviance(p3)/summary(p3)$df.residual,
                  deviance(p4)/summary(p4)$df.residual),

row.names = c('Hour', 'Hour_Month',
              'Hour_Game.Day', 'Hour_Month_Game.Day'))

p_gof
##              Deviance_B GOF_B Dispersion_B Deviance_F GOF_F
## Hour              1112.9389     0      2.243829   1112.9369     0
## Hour_Month         1019.6453     0      2.055737   1019.6337     0
## Hour_Game.Day        971.4373     0      1.958543    971.4301     0
## Hour_Month_Game.Day   933.1296     0      1.881310    933.1241     0
##              Dispersion_B.1
## Hour              2.243824
## Hour_Month         2.068222
## Hour_Game.Day       1.962485
## Hour_Month_Game.Day 1.896594

```

All eight models failed the deviance goodness-of-fit test and showed evidence of being overdispersed. One method of addressing overdispersion is using negative binomial regression.

```

# Bayesian negative binomial regression models adapted from:
# https://georgederpa.github.io/teaching/countModels.html
#
# Accessed on: November 24, 2020
#
#
# Formulas for deviance residuals, log-likelihood, and BIC for
# the negative binomial distribution adapted from formulas
# listed in the following PDF:
# https://ncss-wpengine.netdna-ssl.com/wp-content/themes/ncss/pdf/Procedures/NCSS/Negative\_Binomial\_Regression.pdf
#
# Accessed on: November 29, 2020
#
# Negative binomial models and priors
nb_jags = '
model{
  for(i in 1:n){

    # LIKELIHOODS

```

```

# Model 1 (Count~Hour)
Y1[i] ~ dnegbin(p1[i], r1)
log(lambda1[i]) <- inprod(X1[i,], beta1[])
p1[i] <- r1/(r1+lambda1[i])

# Model 2 (Count~Hour+Month)
Y2[i] ~ dnegbin(p2[i], r2)
log(lambda2[i]) <- inprod(X2[i,], beta2[])
p2[i] <- r2/(r2+lambda2[i])

# Model 3 (Count~Hour+Game.Day)
Y3[i] ~ dnegbin(p3[i], r3)
log(lambda3[i]) <- inprod(X3[i,], beta3[])
p3[i] <- r3/(r3+lambda3[i])

# Model 4 (Count~Hour+Month+Game.Day)
Y4[i] ~ dnegbin(p4[i], r4)
log(lambda4[i]) <- inprod(X4[i,], beta4[])
p4[i] <- r4/(r4+lambda4[i])
}

```

```

# PRIORS

```

```

# Priors for Model 1
r1 ~ dunif(1, 100)
for(i in 1:12){
  beta1[i] ~ dnorm(0, 0.001)
}

```

```

# Priors for Model 2
r2 ~ dunif(1, 100)
for(i in 1:15){
  beta2[i] ~ dnorm(0, 0.001)
}

```

```

# Priors for Model 3
r3 ~ dunif(1, 100)
for(i in 1:13){
  beta3[i] ~ dnorm(0, 0.001)
}

```

```

# Priors for Model 4
r4 ~ dunif(1, 100)
for(i in 1:16){
  beta4[i] ~ dnorm(0, 0.001)
}

```

```

# PREDICTIONS USING THE TEST SET

for(i in 1:m){

  # Model 1 Predictions
  log(lambda1_star[i]) <- inprod(X1_test[i,], beta1[])
  p1_star[i] <- r1/(r1+lambda1_star[i])
  pred1[i] ~ dnegbin(p1_star[i], r1)

  # Model 2 Predictions
  log(lambda2_star[i]) <- inprod(X2_test[i,], beta2[])
  p2_star[i] <- r2/(r2+lambda2_star[i])
  pred2[i] ~ dnegbin(p2_star[i], r2)

  ## Model 3 Predictions
  log(lambda3_star[i]) <- inprod(X3_test[i,], beta3[])
  p3_star[i] <- r3/(r3+lambda3_star[i])
  pred3[i] ~ dnegbin(p3_star[i], r3)

  # Model 4 Predictions
  log(lambda4_star[i]) <- inprod(X4_test[i,], beta4[])
  p4_star[i] <- r4/(r4+lambda4_star[i])
  pred4[i] ~ dnegbin(p4_star[i], r4)
}
}

# Data
nb_data = list(n=nrow(train_X),
               m=nrow(test_X),

               X1=train_X[,1:12],
               X2=train_X[,1:15],
               X3=train_X[,c(1:12, 16)],
               X4=train_X,

               X1_test=test_X[,1:12],
               X2_test=test_X[,1:15],
               X3_test=test_X[,c(1:12, 16)],
               X4_test=test_X,

               Y1=train_Y,
               Y2=train_Y,
               Y3=train_Y,
               Y4=train_Y)

```

```

# Initialize models
nb_models = jags.model(file=textConnection(nb_jags), data=nb_data,
                        inits=list(.RNG.name = 'base::Wichmann-Hill',
                                  .RNG.seed = 7))

## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 2032
##   Unobserved stochastic nodes: 572
##   Total graph size: 41446
##
## Initializing model

# Burn in
update(nb_models, n.iter=1000)

# Sample
nb_outputs = coda.samples(nb_models,
                          variable.names = c('beta1', 'beta2',
                                              'beta3', 'beta4',
                                              'lambda1', 'lambda2',
                                              'lambda3', 'lambda4',
                                              'pred1', 'pred2',
                                              'pred3', 'pred4',
                                              'r1', 'r2',
                                              'r3', 'r4'),
                          n.iter = iter)

# Create 4 frequentist models
nb1 = glm.nb(Count~Hour, data=train_data)
nb2 = glm.nb(Count~Hour+Month, data=train_data)
nb3 = glm.nb(Count~Hour+Game.Day, data=train_data)
nb4 = glm.nb(Count~Hour+Month+Game.Day, data=train_data)

# Extract estimate of each Bayesian node
nb_est = summary(nb_outputs)$statistics[,1]

# Extract Bayesian 95% equi-tailed credible sets
nb_quants = summary(nb_outputs)$quantiles[,c(1,5)]

# Create a table of coefficients with their credible set
nb_beta1 = cbind(Estimate=nb_est[1:12], nb_quants[1:12,])

nb_beta2 = cbind(Estimate=nb_est[13:27], nb_quants[13:27,])

nb_beta3 = cbind(Estimate=nb_est[28:40], nb_quants[28:40,])

```

```

nb_beta4 = cbind(nb_est[41:56], nb_quants[41:56,])

# Rename row names for better interpretation
rownames(nb_beta1) = c('(Intercept)', 'Hour13', 'Hour14', 'Hour15', 'Hour16',
                        'Hour17', 'Hour18', 'Hour19', 'Hour20', 'Hour21',
                        'Hour22', 'Hour23')

rownames(nb_beta2) = c('(Intercept)', 'Hour13', 'Hour14', 'Hour15', 'Hour16',
                        'Hour17', 'Hour18', 'Hour19', 'Hour20', 'Hour21',
                        'Hour22', 'Hour23', 'Month10', 'Month11', 'Month12')

rownames(nb_beta3) = c('(Intercept)', 'Hour13', 'Hour14', 'Hour15', 'Hour16',
                        'Hour17', 'Hour18', 'Hour19', 'Hour20', 'Hour21',
                        'Hour22', 'Hour23', 'Game.Day')

rownames(nb_beta4) = c('(Intercept)', 'Hour13', 'Hour14', 'Hour15', 'Hour16',
                        'Hour17', 'Hour18', 'Hour19', 'Hour20', 'Hour21',
                        'Hour22', 'Hour23', 'Month10', 'Month11', 'Month12',
                        'Game.Day')

```

nb\_beta1

##	Estimate	2.5%	97.5%
## (Intercept)	2.945956561	2.8501333	3.04913398
## Hour13	-0.007702856	-0.1525525	0.13083108
## Hour14	-0.091715678	-0.2436317	0.06088671
## Hour15	0.004687277	-0.1428907	0.14971154
## Hour16	0.021901023	-0.1259354	0.16163502
## Hour17	-0.037374588	-0.1825136	0.10808454
## Hour18	-0.208109587	-0.3583839	-0.06354914
## Hour19	-0.116624234	-0.2753911	0.03742164
## Hour20	-0.311503137	-0.4683523	-0.15678665
## Hour21	-0.309784759	-0.4576187	-0.16190761
## Hour22	-0.294763432	-0.4423637	-0.14737080
## Hour23	-0.290703878	-0.4371789	-0.14452850

nb\_beta2

##	Estimate	2.5%	97.5%
## (Intercept)	3.116355953	3.0053552	3.225979442
## Hour13	-0.017483831	-0.1573635	0.119708926
## Hour14	-0.090901710	-0.2385957	0.060300282
## Hour15	-0.001608053	-0.1484079	0.143059216
## Hour16	0.013509080	-0.1274968	0.151238396
## Hour17	-0.036967859	-0.1762720	0.105216678
## Hour18	-0.222801668	-0.3700270	-0.073179409
## Hour19	-0.146833390	-0.3000294	0.003106644
## Hour20	-0.302444994	-0.4551150	-0.155925716
## Hour21	-0.320962323	-0.4639451	-0.173659792
## Hour22	-0.310708804	-0.4584439	-0.165582601

```
## Hour23      -0.308864734 -0.4587167 -0.159789695
## Month10     -0.158556204 -0.2450749 -0.071856281
## Month11     -0.234613623 -0.3193864 -0.149379886
## Month12     -0.271511606 -0.3534394 -0.187445169
```

nb\_beta3

```
##              Estimate      2.5%      97.5%
## (Intercept)  2.8582438234  2.7619261  2.9580968147
## Hour13      -0.0152006020 -0.1533823  0.1166490073
## Hour14      -0.0859180714 -0.2306254  0.0620140541
## Hour15       0.0001342486 -0.1419478  0.1414802653
## Hour16       0.0272369996 -0.1060675  0.1630312241
## Hour17      -0.0387009109 -0.1800896  0.1007914797
## Hour18      -0.2181378068 -0.3626486 -0.0759936871
## Hour19      -0.1444794149 -0.2952555  0.0005390693
## Hour20      -0.2929693335 -0.4431782 -0.1487326361
## Hour21      -0.3196993937 -0.4615336 -0.1778701986
## Hour22      -0.3070895977 -0.4506681 -0.1594076479
## Hour23      -0.3079636369 -0.4553227 -0.1663394979
## Game.Day     0.2775275766  0.2141724  0.3408365215
```

nb\_beta4

```
##              2.5%      97.5%
## (Intercept)  2.965507092  2.8538003  3.07750122
## Hour13      -0.015633912 -0.1475658  0.11696541
## Hour14      -0.085283369 -0.2262403  0.05831890
## Hour15      -0.006792115 -0.1443675  0.13276158
## Hour16       0.018119686 -0.1088513  0.15011443
## Hour17      -0.038011856 -0.1741755  0.09695838
## Hour18      -0.228967914 -0.3666803 -0.09168618
## Hour19      -0.160650446 -0.3033978 -0.01770115
## Hour20      -0.288928019 -0.4321319 -0.14143492
## Hour21      -0.322875478 -0.4629820 -0.18378739
## Hour22      -0.312880782 -0.4497295 -0.17247506
## Hour23      -0.309369885 -0.4496991 -0.16992463
## Month10     -0.074558344 -0.1637909  0.01248068
## Month11     -0.189735555 -0.2722486 -0.10825826
## Month12     -0.121146744 -0.2118714 -0.03145375
## Game.Day     0.246587793  0.1774711  0.31702508
```

```
# Print GLM coefficients with 95% confidence intervals
cbind(Estimate=coef(nb1), confint(nb1))
```

```
## Waiting for profiling to be done...
```

```
##              Estimate      2.5 %      97.5 %
## (Intercept)  2.944438979  2.8417659  3.04764851
## Hour13      -0.006741599 -0.1508646  0.13735452
## Hour14      -0.090768588 -0.2456736  0.06416725
```

```
## Hour15      0.006397974 -0.1425428  0.15539581
## Hour16      0.023743343 -0.1206710  0.16816026
## Hour17     -0.036141005 -0.1839769  0.11169439
## Hour18     -0.205852054 -0.3567783 -0.05506722
## Hour19     -0.114498960 -0.2710796  0.04210163
## Hour20     -0.310891305 -0.4682076 -0.15381197
## Hour21     -0.310154928 -0.4613573 -0.15920644
## Hour22     -0.293820827 -0.4465428 -0.14132915
## Hour23     -0.288906812 -0.4415247 -0.13651389
```

```
cbind(Estimate=coef(nb2), confint(nb2))
```

```
## Waiting for profiling to be done...
```

```
##           Estimate      2.5 %      97.5 %
## (Intercept)  3.105586057  2.9957839  3.21564715
## Hour13      -0.008244344 -0.1458560  0.12935337
## Hour14      -0.080733634 -0.2288283  0.06731980
## Hour15       0.006349080 -0.1358402  0.14856416
## Hour16       0.021538772 -0.1163044  0.15939233
## Hour17      -0.027018539 -0.1682252  0.11417424
## Hour18      -0.214695680 -0.3591614 -0.07040536
## Hour19      -0.137443567 -0.2873529  0.01238576
## Hour20      -0.292372276 -0.4433239 -0.14172325
## Hour21      -0.311776794 -0.4566403 -0.16719115
## Hour22      -0.301030871 -0.4473778 -0.15495465
## Hour23      -0.298123482 -0.4444990 -0.15201267
## Month10     -0.155997251 -0.2413145 -0.07071792
## Month11     -0.232338454 -0.3171410 -0.14760056
## Month12     -0.269287844 -0.3524635 -0.18618891
```

```
cbind(Estimate=coef(nb3), confint(nb3))
```

```
## Waiting for profiling to be done...
```

```
##           Estimate      2.5 %      97.5 %
## (Intercept)  2.8567477729  2.7590197  2.954484714
## Hour13      -0.0137563763 -0.1477959  0.120273580
## Hour14      -0.0841813176 -0.2285313  0.060085182
## Hour15       0.0008544508 -0.1376598  0.139376033
## Hour16       0.0282255003 -0.1060011  0.162469628
## Hour17      -0.0371556593 -0.1747279  0.100388915
## Hour18      -0.2167359356 -0.3576780 -0.075986754
## Hour19      -0.1423735185 -0.2885383  0.003662098
## Hour20      -0.2917786823 -0.4390566 -0.144845484
## Hour21      -0.3167679053 -0.4582219 -0.175607595
## Hour22      -0.3055505437 -0.4484304 -0.162964459
## Hour23      -0.3070601354 -0.4498717 -0.164539199
## Game.Day1    0.2776555258  0.2149357  0.340417302
```

```
cbind(Estimate=coef(nb4), confint(nb4))
```



```
## Waiting for profiling to be done...
```

```
##           Estimate      2.5 %      97.5 %
## (Intercept)  2.9596691631  2.8469489  3.07231954
## Hour13      -0.0106520122 -0.1418888  0.12058220
## Hour14      -0.0794722721 -0.2208872  0.06182941
## Hour15      -0.0005396158 -0.1361472  0.13506212
## Hour16       0.0234302601 -0.1079662  0.15484596
## Hour17      -0.0321532832 -0.1668639  0.10252454
## Hour18      -0.2222290286 -0.3603803 -0.08428603
## Hour19      -0.1541511793 -0.2974702 -0.01100283
## Hour20      -0.2825127634 -0.4270979 -0.13830094
## Hour21      -0.3162374926 -0.4549750 -0.17780536
## Hour22      -0.3055689943 -0.4456876 -0.16576114
## Hour23      -0.3031745121 -0.4433205 -0.16333277
## Month10     -0.0744623263 -0.1591460  0.01021312
## Month11     -0.1893812587 -0.2713198 -0.10749426
## Month12     -0.1205447133 -0.2104571 -0.03064710
## Game.Day1    0.2467946526  0.1767296  0.31689725
```

Coefficients should match the associated Poisson model (or be approximately the same for the Bayesian models due to the MCMC sampling). However, the variance assumption of Poisson regression is now loosened allowing for a better fit.

```
# Extract fitted values for Bayesian models
```

```
nb_fv1 = nb_est[57:564]
nb_fv2 = nb_est[565:1072]
nb_fv3 = nb_est[1073:1580]
nb_fv4 = nb_est[1581:2088]
```

```
# Extract r for each NB(p,r) distribution
```

```
r1 = nb_est[2601]
r2 = nb_est[2602]
r3 = nb_est[2603]
r4 = nb_est[2404]
```

Negative binomial, NB(p,r), deviance residuals

$$d_i = \text{sign}(Y_i - \lambda_i) \sqrt{2[Y_i \log(Y_i/\lambda_i) - (Y_i + r) \log(\frac{1 + Y_i/r}{1 + \lambda_i/r})]}$$

```
# Calculate deviance residuals
```

```
nb_dev_res = function(fv, r){

  dr = sign(train_Y-fv)*sqrt(2*(train_Y*log(train_Y/fv)
    -(train_Y+r)*log((1+train_Y/r)/(1+fv/r))))

  return(dr)
}
```

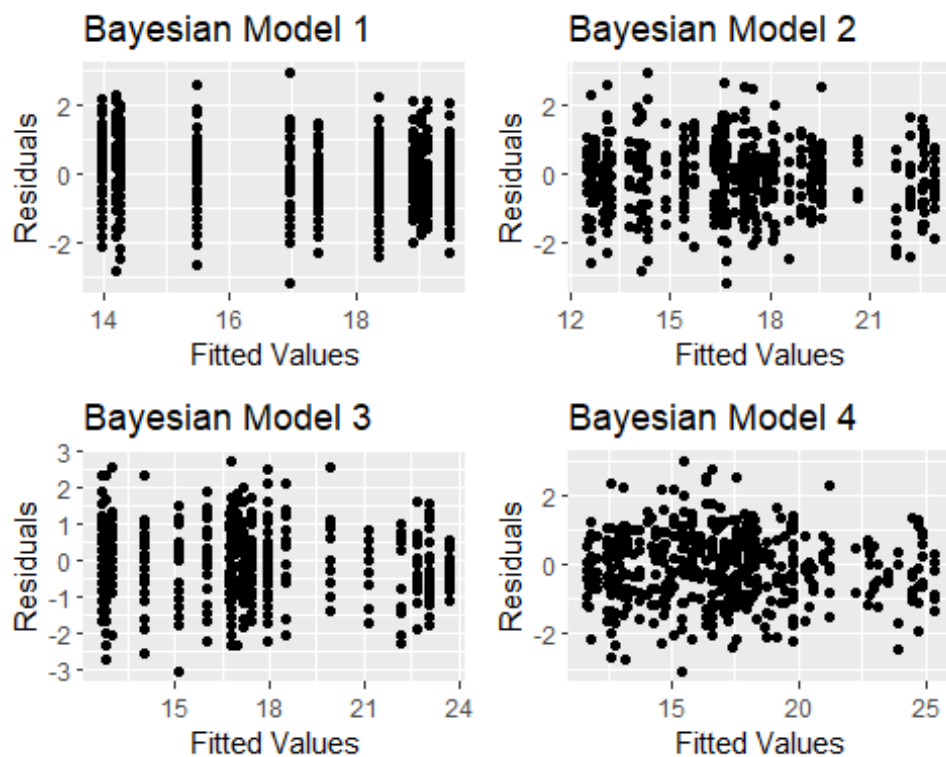
```
# Calculate deviance residuals for Bayesian models
```

```
nb_dr = data.frame(  
  nb_dr1 = nb_dev_res(nb_fv1, r1),  
  nb_dr2 = nb_dev_res(nb_fv2, r2),  
  nb_dr3 = nb_dev_res(nb_fv3, r3),  
  nb_dr4 = nb_dev_res(nb_fv4, r4),  
  row.names = c(1:508))
```

```
# Residual analysis
```

```
# Residuals vs fitted values
```

```
ggarrange(  
  ggplot(nb_dr, aes(x=nb_fv1, y=nb_dr1))+geom_point()+  
    xlab('Fitted Values')+ylab('Residuals')+  
    ggtitle('Bayesian Model 1'),  
  
  ggplot(nb_dr, aes(x=nb_fv2, y=nb_dr2))+geom_point()+  
    xlab('Fitted Values')+ylab('Residuals')+  
    ggtitle('Bayesian Model 2'),  
  
  ggplot(nb_dr, aes(x=nb_fv3, y=nb_dr3))+geom_point()+  
    xlab('Fitted Values')+ylab('Residuals')+  
    ggtitle('Bayesian Model 3'),  
  
  ggplot(nb_dr, aes(x=nb_fv4, y=nb_dr4))+geom_point()+  
    xlab('Fitted Values')+ylab('Residuals')+  
    ggtitle('Bayesian Model 4'))
```



```

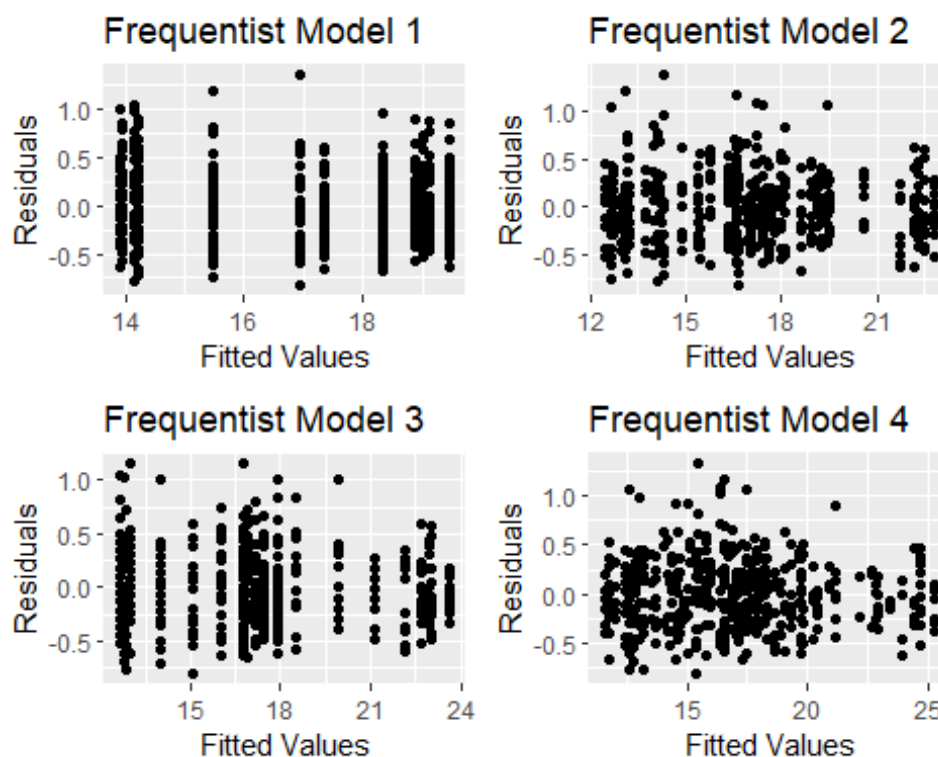
ggarrange(
  ggplot(nb1, aes(nb1$fitted.values, nb1$residuals))+
    geom_point()+xlab('Fitted Values')+
    ylab('Residuals')+ggtitle('Frequentist Model 1'),

  ggplot(nb2, aes(nb2$fitted.values, nb2$residuals))+
    geom_point()+xlab('Fitted Values')+
    ylab('Residuals')+ggtitle('Frequentist Model 2'),

  ggplot(nb3, aes(nb3$fitted.values, nb3$residuals))+
    geom_point()+xlab('Fitted Values')+
    ylab('Residuals')+ggtitle('Frequentist Model 3'),

  ggplot(nb4, aes(nb4$fitted.values, nb4$residuals))+
    geom_point()+xlab('Fitted Values')+
    ylab('Residuals')+ggtitle('Frequentist Model 4'))

```



*# Histogram of the Bayesian residuals*

```

ggarrange(
  ggplot(nb_dr, aes(x=nb_dr1))+geom_histogram()+
    xlab('Deviance Residuals')+ggtitle('Bayesian Model 1'),

  ggplot(nb_dr, aes(x=nb_dr2))+geom_histogram()+
    xlab('Deviance Residuals')+ggtitle('Bayesian Model 2'),

  ggplot(nb_dr, aes(x=nb_dr3))+geom_histogram()+

```


```

xlab('Deviance Residuals')+ggtitle('Bayesian Model 3'),

ggplot(nb_dr, aes(x=nb_dr4))+geom_histogram()+
xlab('Deviance Residuals')+ggtitle('Bayesian Model 4'))

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.

```



```

# Histogram of the frequentist residuals
ggarrange(
  ggplot(nb1, aes(x=nb1$residuals))+geom_histogram()+
  xlab('Deviance Residuals')+ggtitle('Frequentist Model 1'),

  ggplot(nb2, aes(x=nb2$residuals))+geom_histogram()+
  xlab('Deviance Residuals')+ggtitle('Frequentist Model 2'),

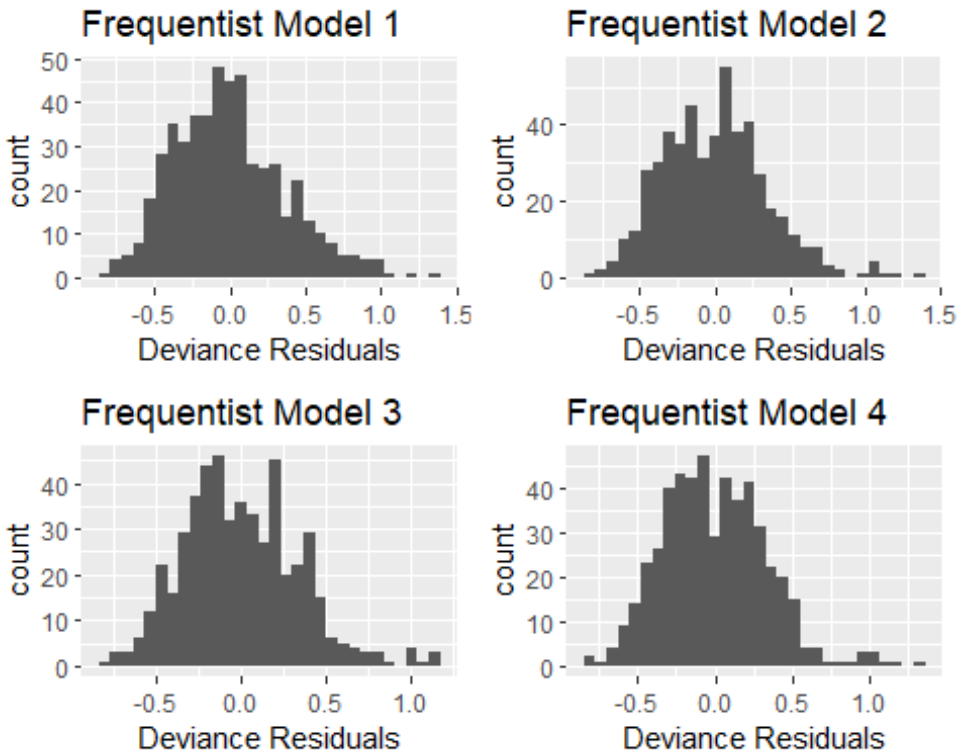
  ggplot(nb3, aes(x=nb3$residuals))+geom_histogram()+
  xlab('Deviance Residuals')+ggtitle('Frequentist Model 3'),

  ggplot(nb4, aes(x=nb4$residuals))+geom_histogram()+
  xlab('Deviance Residuals')+ggtitle('Frequentist Model 4'))

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.

```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
# QQ plot of the Bayesian residuals
ggarrange(ggplot(nb_dr, aes(sample=nb_dr1))+stat_qq()+stat_qq_line()+
  ggtitle('Bayesian Model 1'),

  ggplot(nb_dr, aes(sample=nb_dr2))+stat_qq()+stat_qq_line()+
  ggtitle('Bayesian Model 2'),

  ggplot(nb_dr, aes(sample=nb_dr3))+stat_qq()+stat_qq_line()+
  ggtitle('Bayesian Model 3'),

  ggplot(nb_dr, aes(sample=nb_dr4))+stat_qq()+stat_qq_line()+
  ggtitle('Bayesian Model 4'))
```

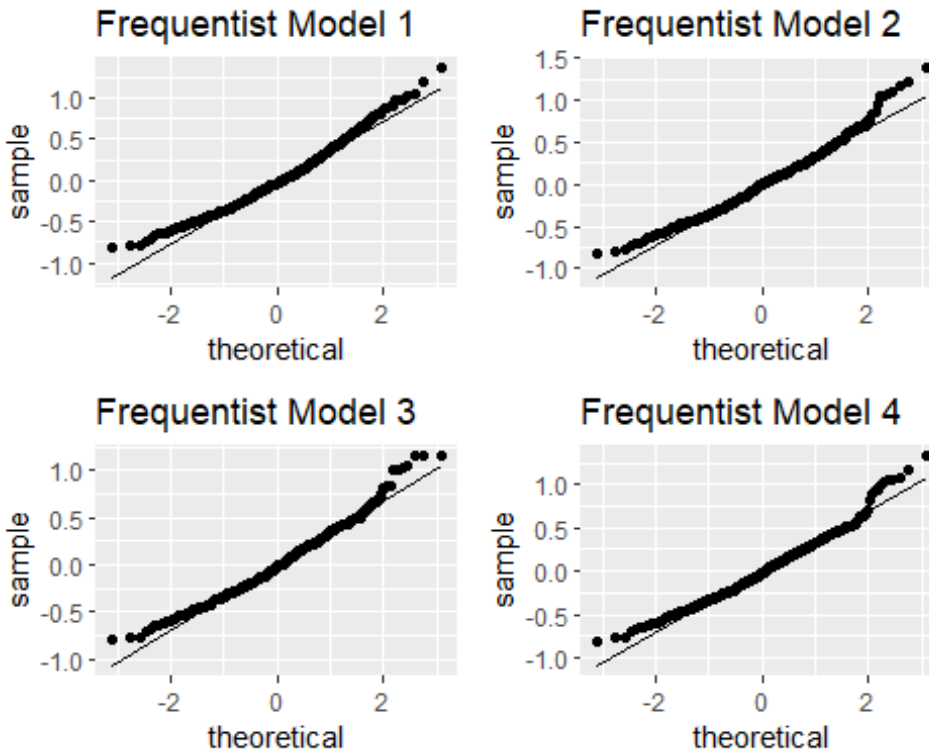


```
# QQ plot of the frequentist residuals
ggarrange(ggplot(nb1, aes(sample=nb1$residuals))+stat_qq()+stat_qq_line()+
  ggtitle('Frequentist Model 1'),

  ggplot(nb2, aes(sample=nb2$residuals))+stat_qq()+stat_qq_line()+
  ggtitle('Frequentist Model 2'),

  ggplot(nb3, aes(sample=nb3$residuals))+stat_qq()+stat_qq_line()+
  ggtitle('Frequentist Model 3'),

  ggplot(nb4, aes(sample=nb4$residuals))+stat_qq()+stat_qq_line()+
  ggtitle('Frequentist Model 4'))
```



Log-likelihood

$LL_i$

$$= \log(\Gamma(Y_i + r)) - \log(\Gamma(r)) - \log(\Gamma(Y_i + 1)) - r \log(1 + \lambda_i/r) - Y_i \log(1 + \lambda_i/r) + Y_i \log(1/r) + Y_i \log(\lambda_i)$$

BIC

$BIC = -2 \sum_{i=1}^n L L_i + p \log(n)$ , where  $p$ =number of predicting variables.

*# Calculate BIC*

```
bayes_bic = function(fv, num_params, r){

  ll = log(gamma(train_Y+r))-log(gamma(r))-log(gamma(train_Y+1))-
    r*log(1+fv/r)-train_Y*log(1+fv/r)+train_Y*log(1/r)+
    train_Y*log(fv)

  return(-2*sum(ll)+num_params*log(508))
}
```

*# Table of BIC values*

```
bic = data.frame(
  BIC_B = c(bayes_bic(nb_fv1, 12, r1),
            bayes_bic(nb_fv2, 15, r2),
            bayes_bic(nb_fv3, 13, r3),
            bayes_bic(nb_fv4, 16, r4)),
```

```

BIC_F = c(AIC(nb1, k=log(508)),
          AIC(nb2, k=log(508)),
          AIC(nb3, k=log(508)),
          AIC(nb4, k=log(508))),

row.names = c('Hour', 'Hour_Month',
              'Hour_Game.Day', 'Hour_Month_Game.Day'))

```

bic

```

##              BIC_B    BIC_F
## Hour          3317.759 3323.949
## Hour_Month     3290.887 3296.984
## Hour_Game.Day   3253.688 3259.874
## Hour_Month_Game.Day 3252.775 3257.635

```

*# Calculate the deviance of each Bayesian model*

```

nb_d1 = sum(nb_dr$nb_dr1^2)
nb_d2 = sum(nb_dr$nb_dr2^2)
nb_d3 = sum(nb_dr$nb_dr3^2)
nb_d4 = sum(nb_dr$nb_dr4^2)

```

*# Summary table (deviance, p-value, dispersion)*

```

nb_gof = data.frame(
  Deviance_B = c(nb_d1, nb_d2, nb_d3, nb_d4),

  GOF_B = c(1-pchisq(nb_d1, summary(nb1)$df.residual),
            1-pchisq(nb_d2, summary(nb2)$df.residual),
            1-pchisq(nb_d3, summary(nb3)$df.residual),
            1-pchisq(nb_d4, summary(nb4)$df.residual)),

  Deviance_F = c(deviance(nb1),
                  deviance(nb2),
                  deviance(nb3),
                  deviance(nb4)),

  GOF_F = c(1-pchisq(deviance(nb1), summary(nb1)$df.residual),
            1-pchisq(deviance(nb2), summary(nb2)$df.residual),
            1-pchisq(deviance(nb3), summary(nb3)$df.residual),
            1-pchisq(deviance(nb4), summary(nb4)$df.residual)),

  row.names = c('Hour', 'Hour_Month',
                'Hour_Game.Day', 'Hour_Month_Game.Day'))

```

nb\_gof

```

##              Deviance_B    GOF_B Deviance_F    GOF_F
## Hour          514.8099 0.2706230   519.7349 0.2227395
## Hour_Month     511.4881 0.2733389   519.0913 0.2010619

```



```
## Hour_Game.Day          514.8000 0.2603701    519.4384 0.2161151
## Hour_Month_Game.Day    483.2372 0.6024672    518.9991 0.1930801
```

#### *# Bayesian predictions*

```
nb_pred1 = nb_est[2089:2216]
nb_pred2 = nb_est[2217:2344]
nb_pred3 = nb_est[2345:2472]
nb_pred4 = nb_est[2473:2600]
```

#### *# GLM predictions*

```
nb1_pred = predict(nb1, test_data, type='response')
nb2_pred = predict(nb2, test_data, type='response')
nb3_pred = predict(nb3, test_data, type='response')
nb4_pred = predict(nb4, test_data, type='response')
```

#### *# Mean Square Prediction Error and Precision Error*

```
mspe = function(pred){
  mspe = mean((pred-test_data$Count)^2)
  return(mspe)
}
```

```
precision = function(pred){
  precision = sum((pred-test_data$Count)^2)/
    sum((test_data$Count-mean(test_data$Count))^2)

  return(precision)
}
```

```
nb_errors = data.frame(
```

```
  MSPE_B = c(mspe(nb_pred1), mspe(nb_pred2),
             mspe(nb_pred3), mspe(nb_pred4)),
```

```
  MSPE_F = c(mspe(nb1_pred), mspe(nb2_pred),
             mspe(nb3_pred), mspe(nb4_pred)),
```

```
  Precision_B = c(precision(nb_pred1), precision(nb_pred2),
                  precision(nb_pred3), precision(nb_pred4)),
```

```
  Precision_F = c(precision(nb1_pred), precision(nb2_pred),
                  precision(nb3_pred), precision(nb4_pred)),
```

```
  row.names = c('Hour', 'Hour_Month',
                'Hour_Game.Day', 'Hour_Month_Game.Day'))
```

```
nb_errors
```

```
##              MSPE_B    MSPE_F Precision_B Precision_F
## Hour          45.17452 45.22076    0.9672709    0.9682612
```

## Hour_Month	38.08498	37.97594	0.8154707	0.8131360
## Hour_Game.Day	38.80083	38.72979	0.8307983	0.8292772
## Hour_Month_Game.Day	36.23542	36.23058	0.7758682	0.7757645